

ELEC0021 - PROGRAMMING II OBJECT-ORIENTED PROGRAMMING

THREADS AND CONCURRENCY

Prof. George Pavlou
Communication and Information Systems Group
Dept. of Electronic and Electrical Engineering
<http://www.ee.ucl.ac.uk/~gpavlou/>
g.pavlou@ucl.ac.uk

Concurrency

- **Concurrency** is a property of running programs to execute computations “in parallel”, overlapped in time
 - These computations share common resources as they have access to the same memory space
- Concurrent use of shared resources is the source of many difficulties
 - Common resources may be accidentally corrupted or get into an inconsistent state when accessed concurrently
- The design of programs to deal with concurrency entails finding reliable techniques for coordinating the execution, data exchange and common memory access among concurrent **execution threads**

Single and Multi-Threaded Programs

- Most programs have a **single thread of execution**, with commands executed sequentially starting from the **main** procedure or method
 - If the program performs many activities, a relatively long activity may delay other activities, e.g. reading new keyboard input, thus making the program unresponsive
 - Also some method calls block until something is available to process, e.g. waiting for a network connection, which means nothing else can be done in the meantime
- With concurrency, the **basic thread** of execution starts from the main method but **additional threads** can be spawned at any time during the program execution, as and when required
 - All the running threads have access to the same memory space, so care is required regarding access to common data
 - If a thread terminates abnormally, e.g. because of an exception, the rest continue and may depend on it -> **important to catch exceptions!**

(Pseudo-) Parallel Execution

- Concurrent tasks realised as threads execute in a “pseudo-parallel” fashion, with the operating system scheduler giving them distinct “CPU time-slices”
 - This is in exactly the same fashion in which the scheduler executes multiple running programs on single-processor systems
- True parallel execution requires threads that exhibit inherent parallelism, i.e. what each thread does is well independent of each other, and a multi-processor system
 - In this case, each processor runs a different thread and threads are swapped in and out in and out of the processors
- In today’s multi-processor systems, typically whole programs rather than threads are delegated to different processors

Concurrent Application Examples

- When downloading and executing audio and video clips over the Internet, the application does not need to wait for the full download before starting playback
 - In this case, synchronisation is required between the two threads
- In a word processor, while a large file is (periodically) saved to disk, the user can continue to word-process “in parallel”
- While a Java program runs, the runtime system may start a thread to reclaim unused memory or otherwise perform “garbage collection” while the program is running
- An application listening for new network connections through a blocking read in the main thread, spawns a new thread to deal a every new connection

Programming Languages and Concurrency

- Initially, support for concurrent programming was only available in an operating system's kernel environment
- **Ada** was the first programming language to support concurrency but it is not a widely-used language
 - It was developed by the US Department of Defence and has been mostly used for military command-and-control systems
- **C** and **C++** use platform-specific libraries for concurrent programming e.g. POSIX threads
 - Porting may be required between different operating systems
- **Java** integrates concurrency support in the core language through **threads of execution**, giving powerful capabilities to the programmer that are available across systems
 - Highly portable in the Java philosophy

Creating and Executing Threads

- The Java thread API has changed in J2SE 5.0, although the old and simpler API is still available
- In J2SE versions < 5.0, threads are supported through the **Thread** class and its methods **start** and **run**
 - **This is the method we will be using**
- In J2SE 5.0, threads are supported through the classes **Executors** and its method **newFixedThreadPool**, **ExecutorService** and its method **execute** and finally the interface **Runnable** and its method **run**
 - This is a more powerful but also more complicated method

The Thread Class

- The Thread class exhibits the key methods **start** and **run**
- Any activity that needs to run as a parallel thread has to be implemented as the method **run** of a specific thread class that extends (i.e. inherits from) class **Thread**
 - E.g. method **run** of class **PrintThread** in the example to follow
- A thread (in addition to the default main program thread) can be started by **creating a new thread object** and calling its **start** method
 - E.g. we are starting three objects of class **PrintThread** in the example to follow
- The Thread class also supports a **sleep** method that will make the thread inactive (“sleeping”) for a number of milliseconds
 - Note: a subclass inherits all the methods of the superclass, e.g. **PrintThread** inherits all the methods of **Thread** such as **start**, **run**, **sleep** and **getName**

A Simple Multi-Threaded Program - 1

```
import java.util.Random;

class PrintThread extends Thread {
    private int sleepTime;

    public PrintThread (String name) {
        // first we call the superclass constructor
        // Thread class keeps a private String name instance variable
        super(name);

        // we now initialise the sleepTime to a random time up to 5 secs
        // this private instance variable will be used in the run method
        Random randGen = new Random();
        sleepTime = randGen.nextInt(5001); // 0 .. 5000 msec
    }

    // continues on the next page
```

A Simple Multi-Threaded Program - 2

// continues from the previous page

```
@override // override Thread's run method (does nothing in Thread)
public void run () {
    try {
        System.out.printf("%s going to sleep for %d msecs\n",
                           getName(), sleepTime);
        Thread.sleep(sleepTime);
        System.out.printf("%s done sleeping, ends\n", getName());
    }
    catch (InterruptedException e) {
        System.err.printf("%s interrupted while sleeping\n", getName());
        e.printStackTrace();
        // printStackTrace is a method of class Throwable from which
        // all exception classes inherit
    }
}

} // end class PrintThread
```

// continues on the next page

A Simple Multi-Threaded Program - 3

// continues from the previous page

```
public class ThreadTester {  
  
    public static void main (String args[])  
    {  
        PrintThread thread1 = new PrintThread("thread1");  
        PrintThread thread2 = new PrintThread("thread2");  
        PrintThread thread3 = new PrintThread("thread3");  
  
        System.out.printf("starting threads 1, 2 & 3\n");  
        thread1.start(); // will eventually invoke run  
        thread2.start(); // ..  
        thread3.start(); // ..  
        System.out.printf("threads scheduled to start, main ends\n");  
    }  
  
} // end class ThreadTester and end of program
```

Some Observations

- The previous example demonstrates threads 1, 2 & 3 started from the main program which sleep for a random time, hence threads 2 and/or 3 may finish before thread 1
 - In fact the program has 4 threads, including the main thread
- It should be noted that the main thread may finish before any of the other three threads start executing:
 - The Thread start method puts a thread in a “ready to execute state” but the actual execution start depends on the scheduler
 - In fact, threads can also be given priorities between 1 (min) and 10 (max), with the default value being 5 (norm); threads with higher priority execute first
- You are advised to run this program in order to see the threads and their execution and ending in action

Synchronisation Requirements

- Multiple threads can have access to common objects, e.g. those created in the main thread and passed to a new thread as constructor arguments
 - This can lead to discrepancies and inconsistencies and subsequently compromise program integrity
- A typical case is based on “read-write” interactions
 - Let’s assume that an instance variable is read, a value is added to it and the instance variable is updated – read and write are implemented through `getValue` and `setValue` methods
 - E.g. Thread 1 reads common object value which is 1000 and adds 200 while Thread 2 does the same and adds 10
 - If both threads do this almost at the same time, Thread 1 will make the value 1200 and then Thread 2 will make it 1010, instead of the correct value that should be 1210 after the two transactions
- The commonly accessed instance variable ends up with a corrupted, i.e. incorrect, value

Synchronisation with Lock

- Java uses **locks** to perform synchronisation. Any object can contain an object that implements the **Lock** interface, typically the **ReentrantLock** one.
- A thread that accesses the containing object needs to do the following to achieve synchronised access:
 - It should first call the Lock's **lock** method; if another thread has the lock, this one will be put in a waiting state
 - It should then call the **getValue** and **setValue** methods; the thread will be allowed to proceed by the runtime system only if the object is not locked by another thread
 - It should finally release the lock by calling the Lock's **unlock** method; **forgetting to do this will block forever any other waiting threads!**

Example Synchronisation with Lock

```
import java.util.concurrent.locks.*;

class SynchronisedReadWrite {
    ReentrantLock accessLock = new ReentrantLock();
    private int value;

    public int getValue () {
        return value;
    }

    public void setValue (int newValue) {
        value = newValue;
    }
}

// ... - srw is a SynchronisedReadWrite object and x the value to add
srw.accessLock.lock();
int srwValue = srw.getValue();
srw.setValue(srwValue + x);
srw.accessLock.unlock();
// ...
```

Summary

- Java threads provide powerful support for concurrency in a platform-independent fashion
- The `Thread` class provides support for concurrency with its `start` and `run` methods
 - Its also supports `getName` and `sleep` methods
- When using threads, locking may be required to lock “critical program regions” that require exclusive access by a single thread